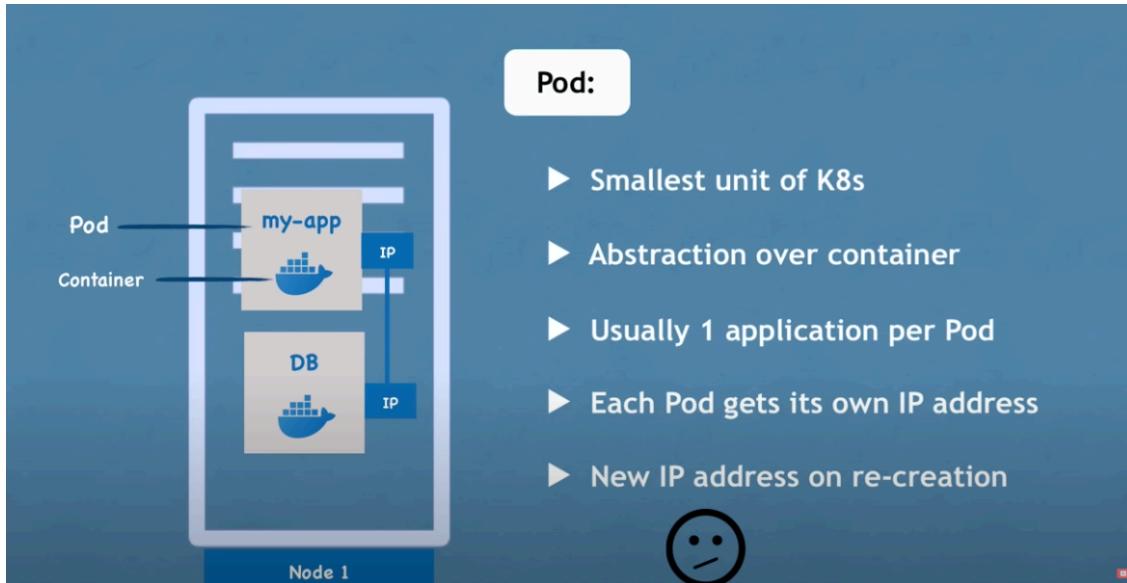
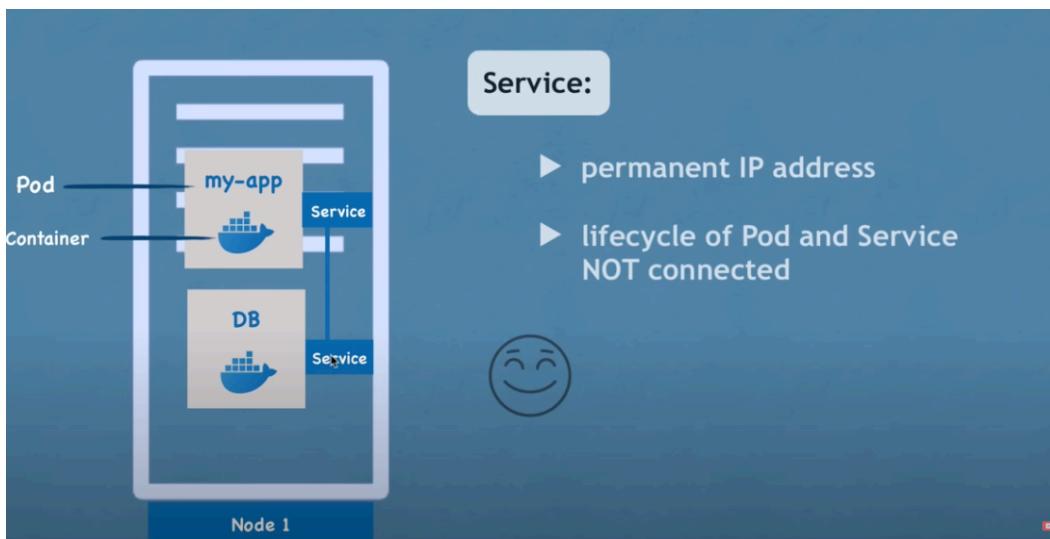


## Kubernetes



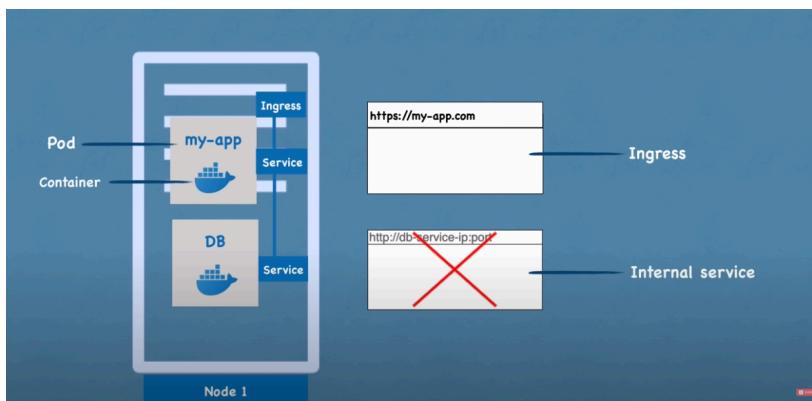
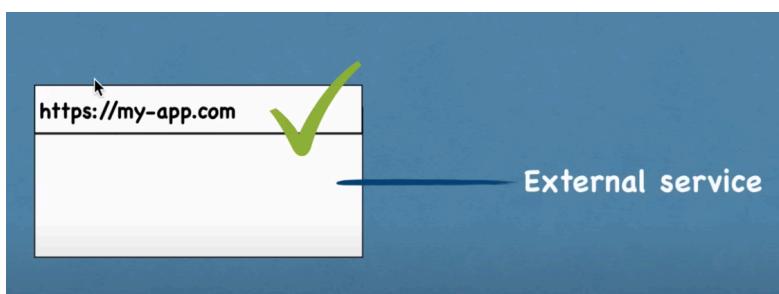
## Why do we need services

If the pods communicate with ip addresses when a pod dies then it could be a problem, because each creation we will have a new ip address. So why we need services



## Why do we need ingress:

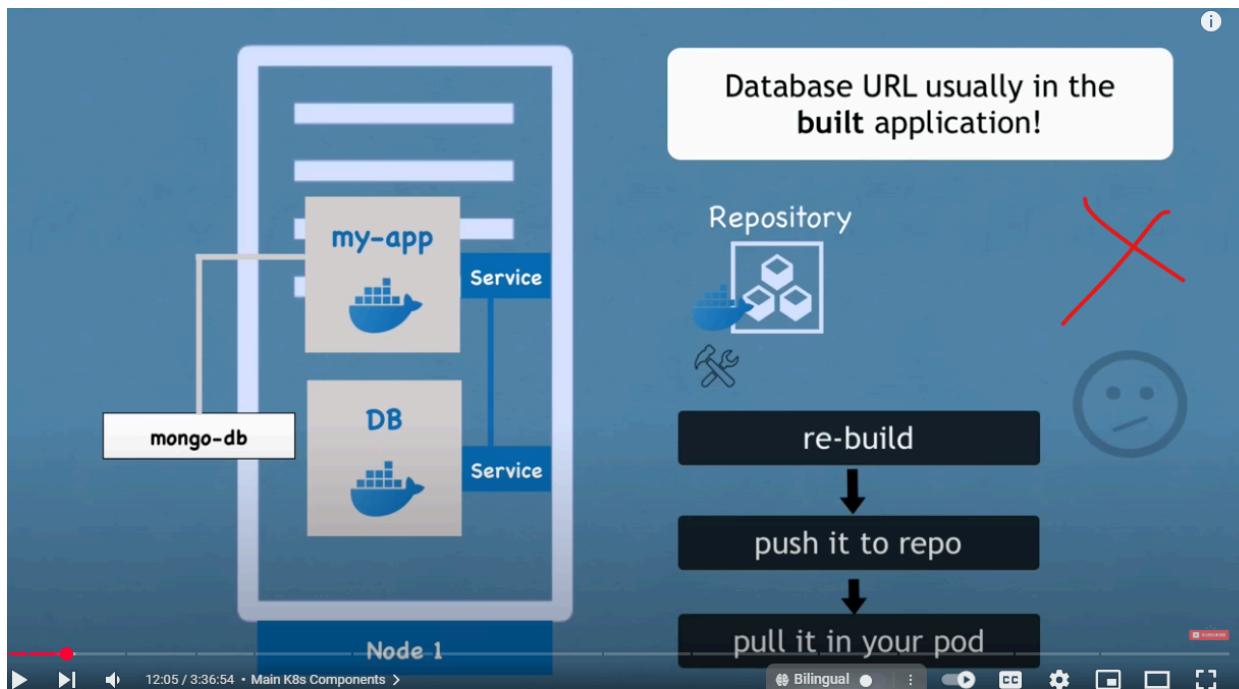
You need **Ingress** in Kubernetes for **external HTTP(S) access** to services within your cluster. Ingress is an API object that manages external access to services, typically HTTP traffic.



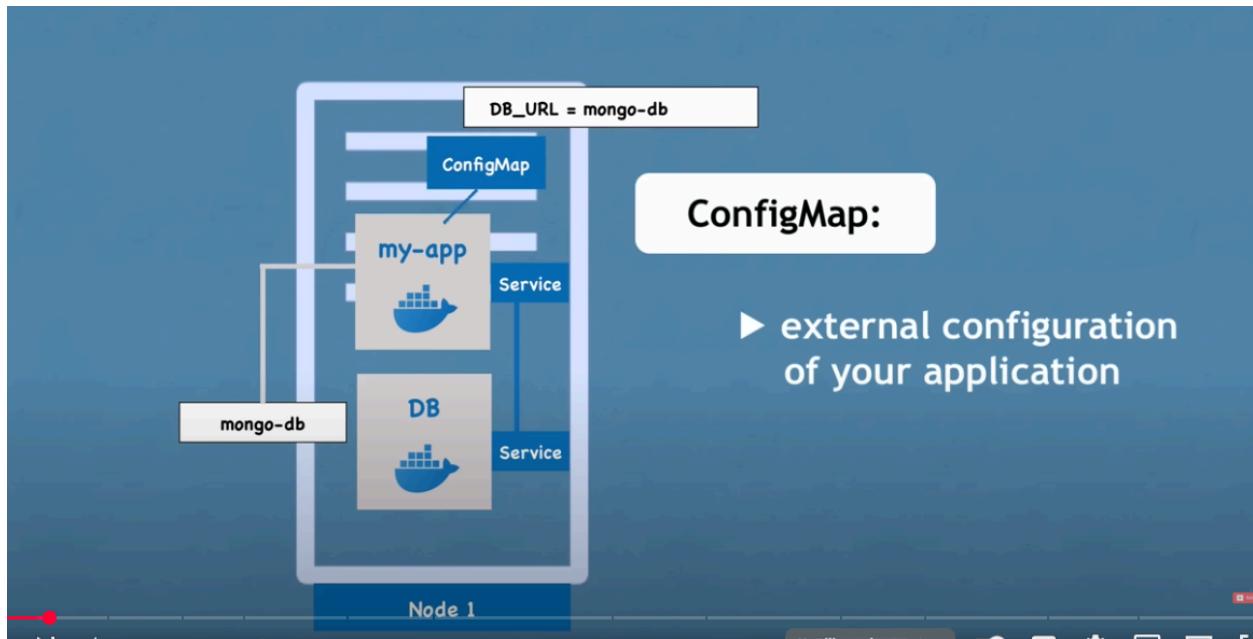
- **Ingress** provides a way to expose services to external traffic.
- It helps with **load balancing**, **SSL termination**, and routing requests to different services based on **URLs or hostnames**.
- **Ingress Controller** is required to process the Ingress rules and route traffic accordingly.

## Why do we need Config Map

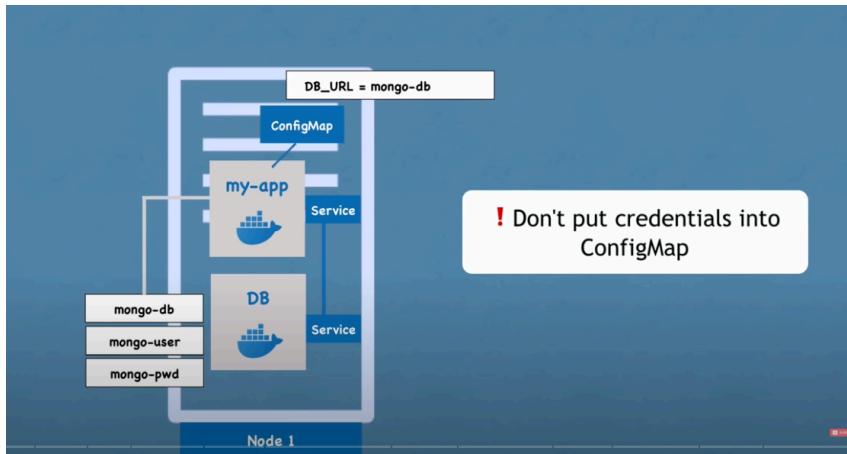
For example if we want to change db url without config map:



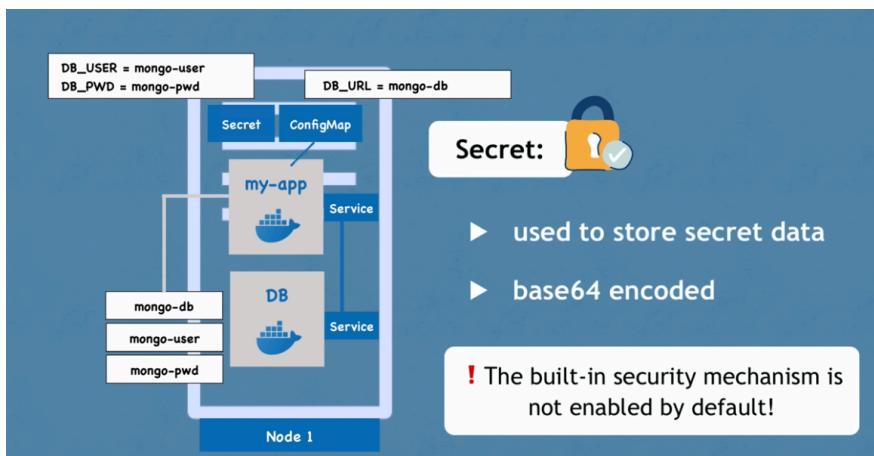
With config map:



Adding for example db credentials into config map is not secure:



Instead:

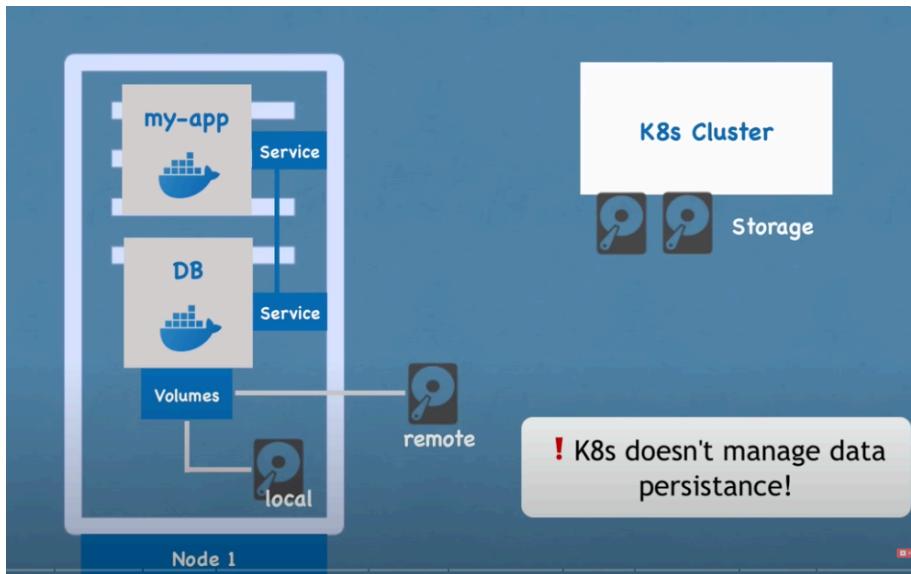
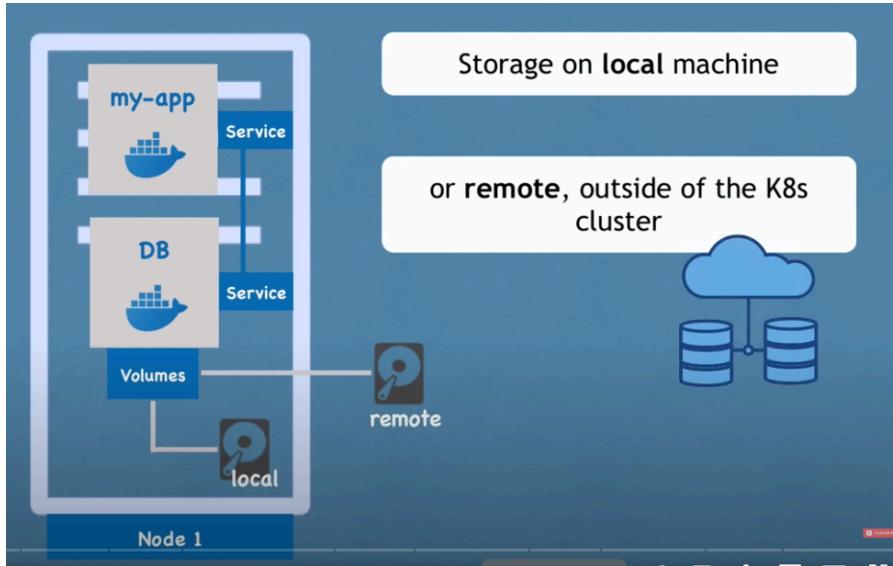


Secrets are **safer than ConfigMaps** because they use [`/etc/secrets/`](#), file permissions, and **etcd encryption**.

Base64 alone is not secure; always enable **etcd encryption** and **RBAC controls**.

## Why do we need volumes?

storing data directly in a Pod is risky because **Pods are ephemeral**. If a Pod dies, any data stored inside it is lost, which can be catastrophic for your application. That's why **persistent storage** is crucial in Kubernetes So why we need volumes



**Kubernetes helps with persistent storage** and makes sure data is available even if Pods fail, by attaching Persistent Volumes to Pods.

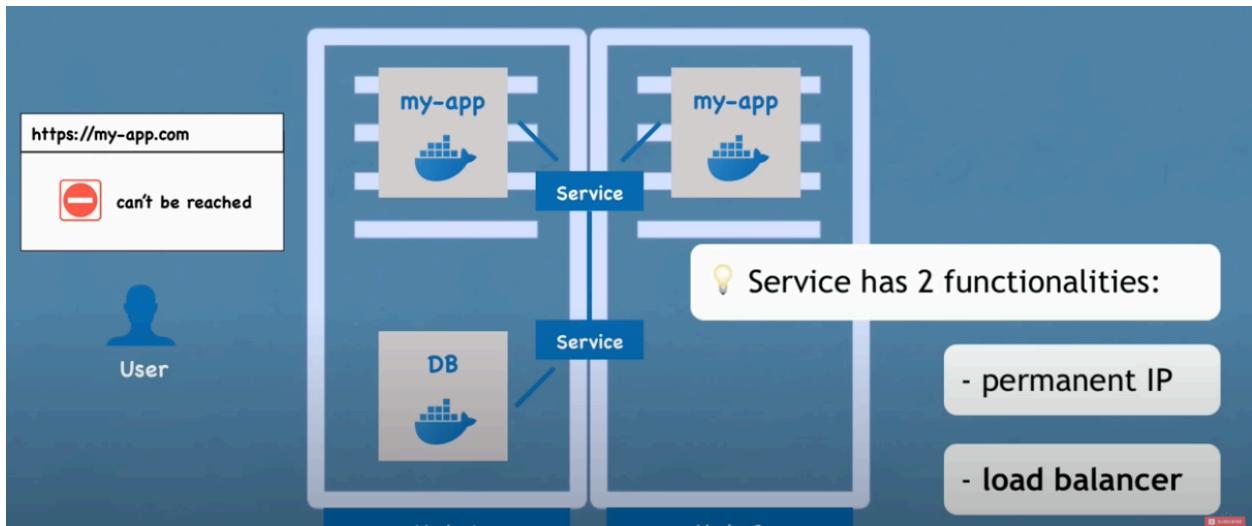
**You are responsible for:**

- **Backups:** Kubernetes doesn't automatically back up your data. You need to set up a backup strategy.
- **Data Replication:** Kubernetes won't replicate your data automatically unless your application (e.g., a database) handles it.

- **High Availability:** For **data redundancy** and **failover** (especially for databases), you have to set it up within your application or use tools like **Replicas**.

### Replica Sets and services

- **ReplicaSet** continuously ensures the correct number of Pods are running.
- The **Service** just keeps routing traffic to whatever Pods are **available** and match its label selector.
- **One Service** can manage multiple replicas. It load-balances traffic across all available Pods.
- Kubernetes ensures that when a **Pod fails**, it replaces it automatically to maintain the desired replica count.
- The Service routes traffic to all Pods that match its label selectors, without needing to know how many replicas exist.



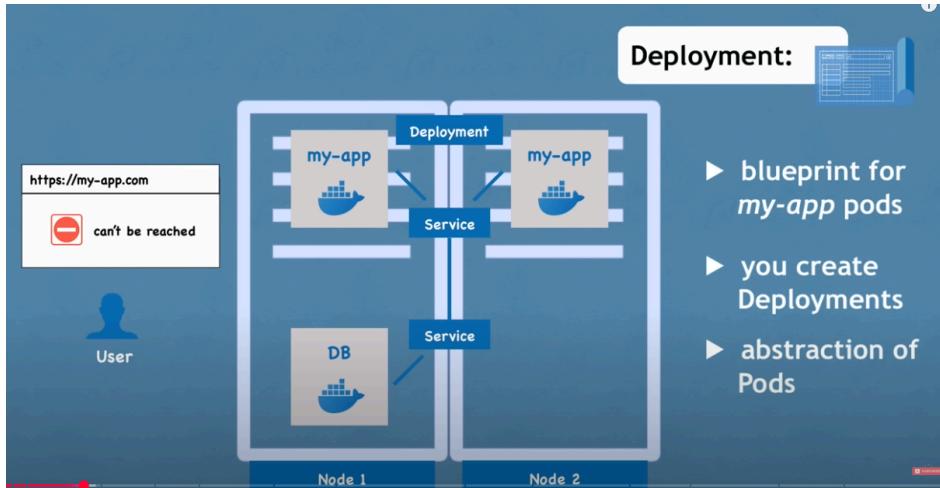
### What Does "Blueprint for an App" Mean?

Blueprint for an app is a template or design that defines how the app should be deployed and run (e.g., using YAML files for Deployments, Services).

- In Kubernetes, this **blueprint** is often represented as **YAML files** (for Deployments, Services, etc.).
- It defines things like:
  - The number of replicas.
  - The container image to use.
  - The resources (CPU, memory) required.
  - Environment variables.
  - Networking details (e.g., ports, services).

A blueprint can be thought of as a **template** that tells Kubernetes how to **deploy** and **manage** the application.

## Deployment



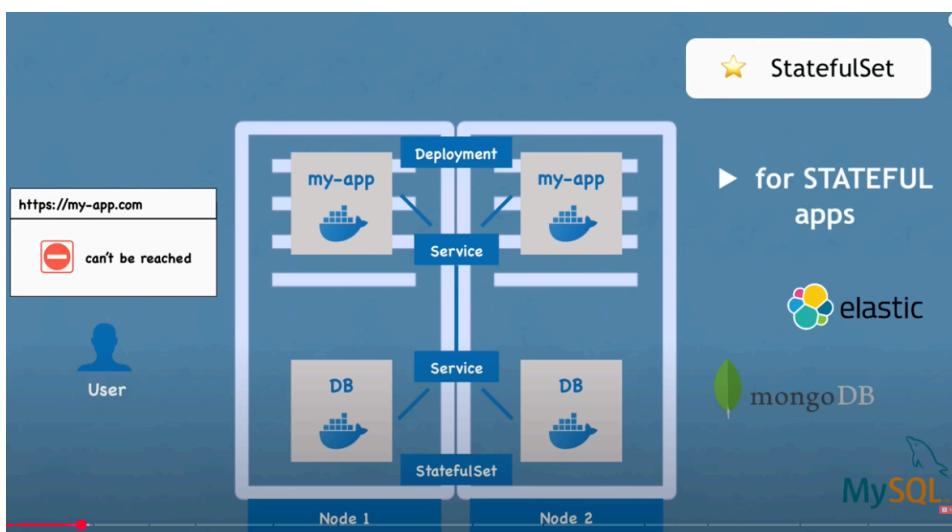
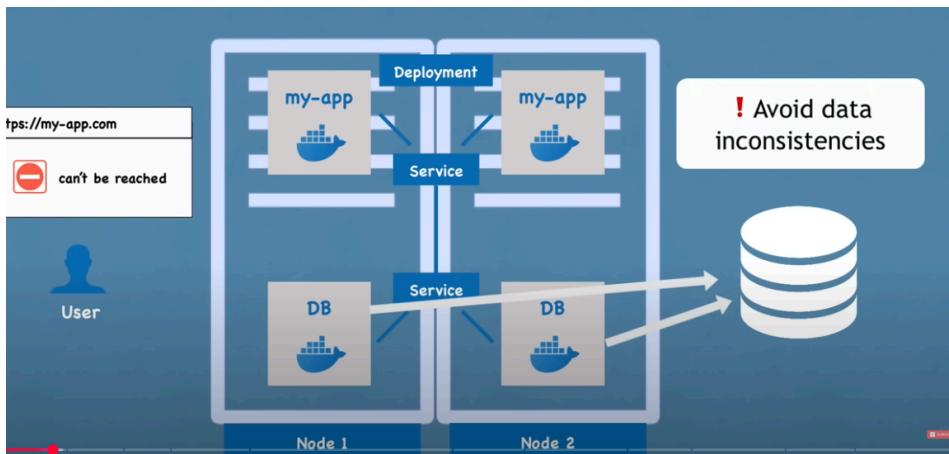
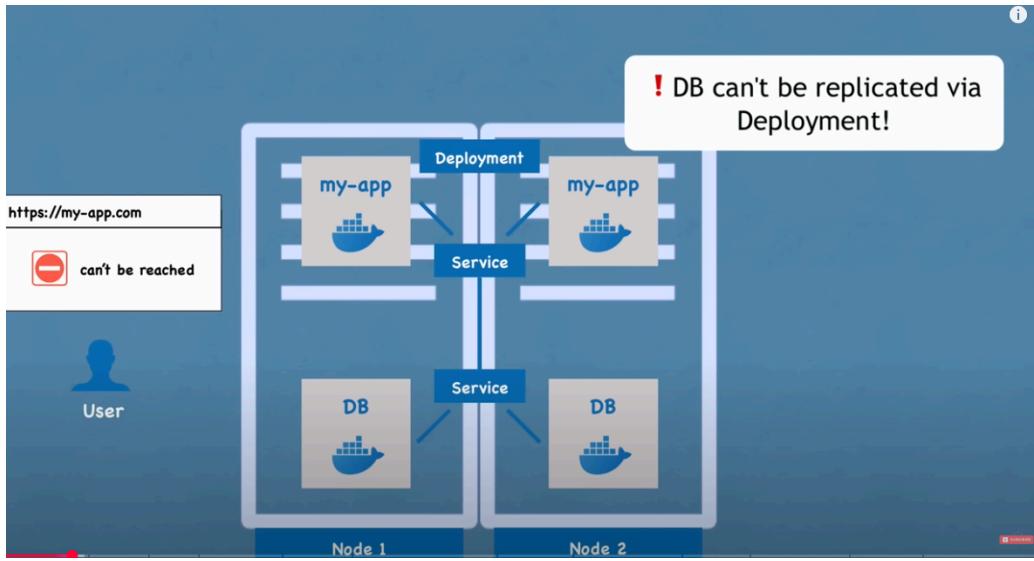
### How Database Replication Works:

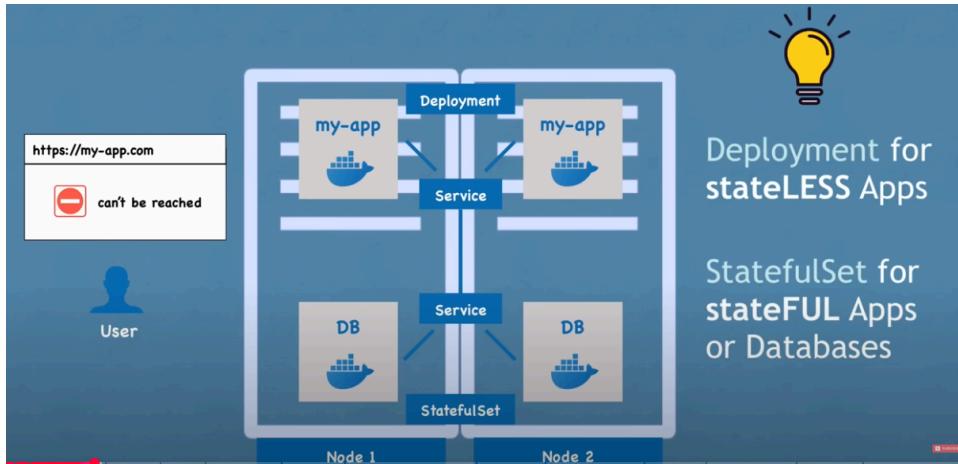
Kubernetes doesn't handle replication for you, but it can manage the **Pods** and **storage**. You need to set up **database replication** inside your database (using tools like MySQL or MongoDB replication).

### What Kubernetes Does:

- Kubernetes manages the **Pods** running your database, using **StatefulSets** instead of Deployments.
- **StatefulSets** ensure each database replica has its own **persistent storage** and stable network identity.

But the actual **replication** (syncing data between database instances) needs to be set up **inside the database** (using MySQL, PostgreSQL, MongoDB replication features, etc.).





- The **primary** database handles **write** operations.
- The **replicas** handle **read** operations.
- **Kubernetes** ensures the databases are running, but **database replication** and **read/write routing** are managed by the **database itself** and your **application** or a **proxy**.

